

Estimating Spoken Vocabulary Size Using Nonparametric Species Richness Methods

Fran Tomazic

March 2026

Abstract

Estimating the total vocabulary of a speaker from a finite recording maps naturally onto the problem of nonparametric species richness estimation, where each distinct word type is treated as a species and each token as an individual. In this thesis, we apply and compare a range of nonparametric estimators, such as Chao1, iChao1, ACE, jackknife, Breakaway, and shape-constrained k-monotone and completely monotone estimators, to transcribed speech data in English, French, and German, using the complete works of Shakespeare as a written replication baseline. The main finding is that estimator choice has a substantially larger effect on the estimated vocabulary size than language: median richness ratios range from around 1.6 for the first-order jackknife to over 5 for Breakaway. Once normalized by the observed vocabulary size. However, these ratios are strikingly consistent across all three languages for any given estimator, suggesting that the gap between what a speaker produces and what they could produce is a relatively universal feature of spoken language. The shape-constrained estimators produce estimates in near-perfect agreement with each other. Furthermore, Breakaway fails on roughly half of all units due to insufficient frequency data, making it unreliable for short spoken samples. In general, moderately flexible estimators, such as Chao1 or ACE offer the best combination of reliability and accuracy for estimating spoken vocabulary.

Contents

1	Introduction	3
2	Framework and Notation	5
2.1	Sampling data structures	5
2.2	Sample coverage	6
2.3	Parametric and nonparametric approaches	7
2.4	Rarefaction, extrapolation, and standardization	8
3	Existing Nonparametric Estimators	10
3.1	Good–Turing frequency estimation	10
3.2	Chao1-type estimators	10
3.3	ACE-type estimators	12
3.4	Jackknife estimators	13
3.5	Summary and limitations	13
4	Beyond Frequency-Count Estimators	15
4.1	Nonparametric mixture models	15
4.2	Frequency-ratio regression (Breakaway)	20
4.3	Shape-constrained estimation	22
4.4	Summary	26
5	Data and Methodology	27
5.1	What is a word?	27
5.2	Data sources	27
5.3	Preprocessing	28
6	Results	29
6.1	Shakespeare replication baseline	29
6.2	Corpus and language overview	31
6.3	Per-language estimator results	33
6.4	Completely monotone and k-monotone model selection	33
6.5	Rarefaction and extrapolation	37
6.6	Cross-language comparison	37
6.7	Estimator agreement and disagreement	39
6.8	Summary	41
7	Conclusion	42
A	Parametric mixture models	43
B	Data Sources	43

1 Introduction

The problem of estimating species richness has a long history in multiple disciplines. It is the total number of distinct species present in an assemblage, including those not observed in any sample. The 1943 paper by Fisher, Corbet, and Williams [10] laid the mathematical foundation for statistical approaches to this problem. Around the same time, Good and Turing developed frequency estimation methods as part of their effort to crack German ciphers during World War II [12, 13]. Since then, the “species problem” has appeared in ecology (estimating butterfly species on an island), biology (counting microbial taxa), genetics (enumerating alleles), software engineering (finding undiscovered bugs), and linguistics (determining how many words an author knew) [5].

In nearly all of these settings, counting the true richness is almost impossible. A sample will almost always miss a fraction of the true existing species, and the simple count of observed species S_{obs} systematically underestimates the true richness S . The magnitude of this negative bias depends on the heterogeneity of the underlying abundance distribution: the more unequal the species abundances, the more rare species the sample will fail to detect. This makes the gap between S_{obs} and S the central statistical challenge of the field.

Our main question throughout this paper is: *how can we estimate the total spoken vocabulary of a person from a finite recording of their speech?* In this setting, each distinct word corresponds to a species, and each token (word occurrence) corresponds to an individual. Later on, we will define more precisely what we mean by a distinct word. As we can see, spoken vocabulary falls right into the regime where species richness estimation is most difficult: word frequencies are highly heterogeneous. They are governed by Zipf’s law, which states that the frequency of a word is approximately inversely proportional to its rank [1]. What is the rank in this context? Suppose that you order all the words of a text by their frequencies, then the highest frequency word has rank 1, second highest has rank 2, etc... and by Zipf’s law formula we get that $f(1) \propto 1$ and $f(2) \propto 1/2$ thus the rank 1 word appears twice as often as the rank 2 word. This comes from the fact that a small number of words (*the, is, a and*) represent a large share of all tokens, while the majority of distinct words appear only once or twice in any finite sample. Therefore, the number of words observed exactly once or twice, commonly referred to as singletons and doubletons, is always large.

A person’s *active* spoken vocabulary (with how many different words they speak on a daily basis) is considerably smaller than their *passive* vocabulary (the words they can understand). Research in applied linguistics suggests that comprehension of spoken texts requires knowledge of roughly 6 000 to 7 000 word families, whereas written texts require 8 000 to 9 000. But what a person *produces* in conversation is a much smaller, context-dependent subset of their active vocabulary. Estimating the size of this active vocabulary from a finite speech sample is therefore a nontrivial statistical problem, and one that is well suited to the tools of nonparametric species richness estimation. We will also

compare estimates derived from spoken data with those from written text, where we expect greater lexical diversity due to the wider range of registers and the time available for word choice.

In this thesis, we proceed as follows. In Section 2 we establish the mathematical framework and notation for species richness estimation, distinguishing abundance data from incidence data and defining sample coverage. Section 3 reviews classical nonparametric estimators such as Good-Turing, Chao1, iChao1, ACE, and the jackknife, discussing the assumptions, strengths, and limitations of each. Section 4 presents the main theoretical contribution: shape-constrained estimation via k -monotone and completely monotone distributions. We explain why using the completely monotone model is a natural next step when the degree of monotonicity k is large. Section 5 describes our data sources: the Shakespeare word counts from Spevack [18] for replication, and several other spoken language corpora. Section 6 presents the results, including estimator comparisons, rarefaction, and extrapolation curves. We discuss the findings and conclude in Section 7.

2 Framework and Notation

Let us define our framework and notation. Consider an assemblage consisting of N total individuals, which is a finite set, each belonging to one of S distinct species. Let N_i denote the number of individuals of the i^{th} species in the assemblage, with $i = 1, 2, \dots, S$ and $N_i > 0$. The total population size is

$$N = \sum_{i=1}^S N_i.$$

The relative abundance of species i is $p_i = N_i/N$, with $\sum_{i=1}^S p_i = 1$. The quantities N , S , N_i , and p_i represent the true and unknown parameters of the assemblage.

Assume that we draw a random sample of n individuals from the assemblage, thus having $n < N$. In this sample, a total of S_{obs} distinct species are observed, we call it the *reference sample*. Let X_i be the number of individuals of the i^{th} species observed in this sample. Only species with $X_i > 0$ are observable, and $\sum_{i=1}^S X_i = n$. We denote by f_k the *abundance frequency count*: the number of species represented by exactly k individuals in the reference sample for $k = 1, 2, \dots, n$. That is,

$$f_k = \#\{i : X_i = k\}.$$

The sample size can then be expressed in terms of the abundance frequency counts as

$$n = \sum_{i=1}^S X_i = \sum_{k \geq 1} k f_k.$$

We call f_1 the number of *singletons* (species observed exactly once) and f_2 the number of *doubletons* (species observed exactly twice). Singletons and doubletons play a central role in many estimation methods, as we will see in Section 3. More generally, most of the information on undetected species is contained in the lower-order frequency counts $f_1, f_2, \dots, f_r$ for small r and the quantity we wish to estimate is f_0 : the number of species present in the assemblage but absent from the reference sample. The true species richness decomposes as

$$S = S_{\text{obs}} + f_0. \tag{1}$$

Since S_{obs} is observed directly, the entire estimation problem reduces to estimating f_0 .

2.1 Sampling data structures

There are two principal types of sampling data: abundance and incidence data.

Abundance data. In abundance sampling, the sampling unit is an individual. In the discrete-type variant, a fixed number n of individuals is drawn from a study area. In the continuous-type variant, the sampling effort is measured on a continuous scale such as time, area, or volume. For our application, we record a speaker for a fixed period of time, and each word said constitutes an individual drawn from the speaker’s vocabulary. The resulting data are discrete-type abundance data.

Incidence data. In incidence sampling, the sampling unit is not an individual, but a spatial or temporal unit (a quadrat, a time window, etc.). For each of the T sampling units, one records only the detection or non-detection of each species, yielding an incidence matrix $[W_{ij}]$ and incidence-based frequency counts Q_k ($k = 1, \dots, T$) i.e., the number of species detected in exactly k sampling units. Therefore, Q_T refers to the number of species detected in all sampling units. The number of all observed species is given by $\sum_{k=1}^T Q_k$ and since our data consist of individual words with their frequencies, we work mostly with abundance data throughout this thesis. For a detailed treatment of incidence-based methods, we refer to Chao and Chiu [5].

2.2 Sample coverage

An important concept for assessing how complete a sample is is the notion of *sample coverage*. Introduced by Good and Turing [12], it estimates the proportion of the total number of individuals in the assemblage that belong to species represented in the sample. Turing’s estimator is

$$\hat{C} = 1 - \frac{f_1}{n}. \quad (2)$$

Intuitively, f_1/n estimates the probability that the next sampled individual belongs to a species not yet observed, subtracting it from 1 gives the coverage. A coverage close to 1 indicates that most of the assemblage’s individuals belong to species already detected, whereas a low coverage signals that many species remain unobserved. The same number of singletons can result from a small homogeneous vocabulary sampled sparsely or from a vast heterogeneous one sampled thoroughly, so $\hat{C}$ cannot distinguish between these two fundamentally different scenarios; therefore, Chao and Jost [6] proposed a bias-corrected version:

$$\hat{C} = 1 - \frac{f_1}{n} \cdot \frac{(n-1)f_1}{(n-1)f_1 + 2f_2}. \quad (3)$$

This adjustment reduces the bias of the basic Turing estimator, particularly when singletons dominate the sample (i.e. when f_1 is large relative to n). Sample coverage will later play an important role, both as a diagnostic of the completeness of our speech samples and as the basis for coverage-based standardization when comparing samples of different sizes.

Figure 1 illustrates the scale of the coverage values in concrete examples. A high value of $\hat{C}$ indicates that most of the probability mass belongs to the word types already seen in the sample. A lower value indicates that a substantial part of the vocabulary is still represented only through rare types, especially singletons.

Examples of Sample Coverage in the Thesis Corpora

Example	Tokens n	Observed types S_{obs}	f_1	f_2	Turing $\hat{C}$
IMSDb character 1	650	345	249	49	0.6169
IMSDb character 2	552	249	155	40	0.7192
BNC speaker 1	558	217	137	31	0.7545
SBCorpus speaker 1	1,483	558	346	99	0.7667
SBCorpus speaker 2	1,443	426	257	62	0.8219
IMSDb character 3	3,422	967	580	146	0.8305
BNC speaker 2	1,598	431	213	69	0.8667
SBCorpus speaker 3	2,615	554	284	92	0.8914
Hamlet	30,095	4,668	2,768	704	0.9080
BNC speaker 3	15,351	2,013	1,016	359	0.9338
Full Shakespeare corpus	890,426	25,700	9,704	3,562	0.9891

Figure 1: Examples of Turing sample coverage for several texts and speakers used in this thesis.

2.3 Parametric and nonparametric approaches

A parametric model assumes that the abundance distribution of the species belongs to a known family of distributions, fully characterized by a finite parameter vector θ . Common choices include log-normal, inverse-Gaussian, and Zipf-Mandelbrot law. Although such models can produce efficient estimators when the assumed form is well specified, simulations have shown that they perform poorly when not [5]. Moreover, two parametric models may fit the observed data equally well, but produce widely different estimates of S , making the choice of distributional family a source of considerable uncertainty (as illustrated in Figure 3).

A nonparametric model makes no assumption about the functional form of the underlying abundance distribution. Instead, the number of undetected species is estimated directly from the lower-order frequency counts $f_1, f_2, \dots, f_r$ for some small species r . This approach is more robust across applications and permits meaningful comparisons between assemblages since it does not require that different assemblages follow the same distributional family.

2.4 Rarefaction, extrapolation, and standardization

When comparing species richness across multiple assemblages (for example, the vocabulary of two different speakers, or spoken versus written language from the same individual), a direct comparison of S_{obs} is misleading if sample sizes differ, since larger samples will generally detect more species. Two standardization approaches are commonly used.

Sample-size-based rarefaction and extrapolation. Rarefaction asks: if we had drawn only $m < n$ individuals, how many species would we expect to observe? By subsampling from the reference sample, one constructs a *rarefaction curve* that shows the expected richness as a function of the sample size. In contrast, extrapolation projects the curve beyond the observed sample size, predicting how many additional species would be detected if the sampling effort were increased [6].

Figure 2 shows why the observed raw richness is not directly comparable between samples of different sizes. Even for a fixed text, the expected number of observed word types changes substantially with the number of sampled tokens. Rarefaction and extrapolation therefore provide a common sampling-effort scale on which different assemblages can be compared.




Figure 2: Sample-size-based rarefaction and extrapolation for one example text. The solid curve gives the expected observed richness for smaller sample sizes $m < n$, the point marks the observed sample (n, S_{obs}) , and the dashed curve extrapolates richness up to twice the observed sample size.

Coverage-based standardization. An alternative, proposed by Chao and Jost [6], standardizes samples not by size, but by coverage. The rationale is that two samples of equal size may represent very different levels of completeness: a sample of $n = 5\,000$ tokens from a speaker with homogeneous word usage may achieve near-complete coverage, while the same n from a speaker with a highly heterogeneous vocabulary may miss a substantial fraction of species. By comparing richness with equal coverage, one obtains a fairer comparison that accounts for the difficulty of each assemblage. We will use both approaches when comparing our spoken and written data in Section 6.

3 Existing Nonparametric Estimators

In this section, we review the principal nonparametric estimators of species richness. All estimate the number of undetected species f_0 using only the lower-order frequency counts, without assuming any particular distributional form for the abundances of species.

3.1 Good–Turing frequency estimation

The Good–Turing frequency estimator, developed by I. J. Good based on ideas from Alan Turing, originated in their wartime work on cracking German ciphers [12, 13]. Good initially modeled the species frequencies using a multinomial distribution, but found this approach inaccurate for rare species. The key insight of the Good–Turing estimation is to adjust the observed frequency of a species: for a species observed k times, its estimated probability is

$$\hat{p}_k = \begin{cases} \frac{(k+1)f_{k+1}}{n f_k}, & k > 0, \\ \frac{f_1}{n}, & k = 0. \end{cases} \quad (4)$$

where n is the total sample size. $\hat{p}_0$ is *not* the count of undetected species f_0 ; rather, it estimates the fraction of the assemblage’s individuals that belong to species absent from the sample. It is the complement of the sample coverage $\hat{C} = 1 - f_1/n$ defined in Section 2.2.

Simple Good–Turing smoothing. A practical difficulty with Equation (4) is that for large k , the frequency counts f_k become noisy and sparse, making the ratio f_{k+1}/f_k unstable. The Simple Good–Turing (SGT) algorithm addresses this by replacing the raw counts with smoothed values $S(f_k)$ obtained from a log-linear regression of $\log f_k$ in $\log k$ [11]:

$$p_k = \frac{(k+1) \cdot S(f_{k+1})}{n \cdot S(f_k)}. \quad (5)$$

For small values of k , where the counts are reliable, one uses $S(f_k) = f_k$ directly, the smoothed values are used only in the sparse tail. We will not go into further detail on the regression procedure here as the SGT algorithm is not the main focus of our analysis.

3.2 Chao1-type estimators

The Chao1 estimator, introduced by Chao [4], is the most widely used non-parametric lower bound for species richness. Its derivation applies the Cauchy-Schwarz inequality to the expected frequency counts, which guaranties that the result is a lower bound for S regardless of the true abundance distribution. Now we present the derivation.

Derivation. Assume that each species i has a detection rate λ_i drawn from some distribution G , and that, conditional on λ_i , the species count follows a distribution $\text{Poisson}(\lambda_i)$ (this framework is fully developed in Section 4.1). The expected frequency counts then satisfy the following.

$$\mathbb{E}[f_j] = S \int_0^\infty \frac{\lambda^j e^{-\lambda}}{j!} dG(\lambda), \quad j = 0, 1, 2, \dots$$

In particular,

$$\begin{aligned} \mathbb{E}[f_0] &= S \int_0^\infty e^{-\lambda} dG(\lambda), \\ \mathbb{E}[f_1] &= S \int_0^\infty \lambda e^{-\lambda} dG(\lambda), \\ \mathbb{E}[f_2] &= \frac{S}{2} \int_0^\infty \lambda^2 e^{-\lambda} dG(\lambda), \\ \mathbb{E}[f_3] &= \frac{S}{6} \int_0^\infty \lambda^3 e^{-\lambda} dG(\lambda). \end{aligned}$$

Define the probability measure $dH(\lambda) = e^{-\lambda} dG(\lambda) / \int e^{-\lambda} dG(\lambda)$, so that expectations can be rewritten as

$$\mathbb{E}[f_1] = \mathbb{E}[f_0] \int_0^\infty \lambda dH(\lambda), \quad \mathbb{E}[f_2] = \frac{\mathbb{E}[f_0]}{2} \int_0^\infty \lambda^2 dH(\lambda).$$

By the Cauchy–Schwarz inequality, $(\int \lambda dH)^2 \leq \int \lambda^2 dH$, which gives

$$(\mathbb{E}[f_1])^2 \leq 2 \mathbb{E}[f_0] \mathbb{E}[f_2].$$

Rearranging yields a lower bound on the number of undetected species:

$$\mathbb{E}[f_0] \geq \frac{(\mathbb{E}[f_1])^2}{2 \mathbb{E}[f_2]}. \quad (6)$$

Replacing $\mathbb{E}[f_j]$ by observed f_j and adding S_{obs} gives the Chao1 estimator (with a finite-sample bias correction $\frac{n-1}{n}$):

$$\hat{S}_{\text{Chao1}} = \begin{cases} S_{\text{obs}} + \frac{n-1}{n} \cdot \frac{f_1^2}{2f_2}, & \text{if } f_2 > 0, \\ S_{\text{obs}} + \frac{n-1}{n} \cdot \frac{f_1(f_1-1)}{2}, & \text{if } f_2 = 0. \end{cases} \quad (7)$$

For large sample sizes n , the correction factor $(n-1)/n$ is negligible, and the estimator simplifies to $\hat{S}_{\text{Chao1}} \approx S_{\text{obs}} + f_1^2/(2f_2)$.

The Cauchy–Schwarz inequality is sharp: equality holds if and only if λ is constant under H , i.e., all species have the same detection rate. This means that the Chao1 estimator achieves the true value only in the homogeneous case where all species are equally detectable. In practice, this condition is rarely met for linguistic data, where word frequencies are highly heterogeneous. When heterogeneity is high, the true f_0 can substantially exceed the Chao1 estimate, meaning the estimator underestimates the true richness [5, 17].

The iChao1 estimator. Chiu et al. [9] derived an improved lower bound by extending the Cauchy-Schwarz argument to incorporate tripletons f_3 and quadrupletons f_4 :

$$\hat{S}_{i\text{Chao1}} = \hat{S}_{\text{Chao1}} + \frac{n-3}{n} \cdot \frac{f_3}{4f_4} \times \max\left(f_1 - \frac{n-3}{n-1} \cdot \frac{f_2 f_3}{2f_4}, 0\right). \quad (8)$$

For large n , this simplifies to $\hat{S}_{i\text{Chao1}} \approx \hat{S}_{\text{Chao1}} + \frac{f_3}{4f_4} \max\left(f_1 - \frac{f_2 f_3}{2f_4}, 0\right)$. By incorporating higher-order frequency counts, iChao1 tightens the lower bound when there is additional information in f_3 and f_4 beyond what f_1 and f_2 alone provide.

3.3 ACE-type estimators

The Abundance-based Coverage Estimator (ACE), developed by Chao and Lee [7], expands the idea of sample coverage by incorporating a measure of heterogeneity among rare species. The central idea is that abundant species are almost certain to appear in any sample and therefore do not carry information on the undetected fraction f_0 . Only rare species, whose detection is uncertain, are informative. ACE formalizes this by splitting the observed species into two groups using a cutoff value k (conventionally $k = 10$, following Chao and Lee [7]):

$$\begin{aligned} S_{\text{abun}} &= \sum_{i>k} f_i, & S_{\text{rare}} &= \sum_{i=1}^k f_i, \\ n_{\text{rare}} &= \sum_{i=1}^k i f_i, & \hat{C}_{\text{rare}} &= 1 - \frac{f_1}{n_{\text{rare}}}, \end{aligned}$$

where n_{rare} is the sample size restricted to the rare group and $\hat{C}_{\text{rare}}$ is Turing's coverage applied to that group.

The key quantity is the squared coefficient of variation (CV) of the species abundances:

$$\gamma^2 = \frac{S^{-1} \cdot \sum_{i=1}^S (p_i - \bar{p})^2}{\bar{p}^2}, \quad (9)$$

where $\bar{p} = 1/S$. This characterizes the degree of heterogeneity between species abundances: $\gamma^2 = 0$ if and only if all species have equal abundances (homogeneous assemblage), and higher values of γ^2 indicate greater heterogeneity. The ACE estimator is then given by the following:

$$\hat{S}_{\text{ACE}} = S_{\text{abun}} + \frac{S_{\text{rare}}}{\hat{C}_{\text{rare}}} + \frac{f_1}{\hat{C}_{\text{rare}}} \cdot \hat{\gamma}_{\text{rare}}^2, \quad (10)$$

with the estimated squared CV of the rare group given by

$$\hat{\gamma}_{\text{rare}}^2 = \max \left\{ \frac{S_{\text{rare}}}{\hat{C}_{\text{rare}}} \cdot \frac{\sum_{i=1}^k i(i-1) f_i}{\left(\sum_{i=1}^k i f_i\right) \left(\sum_{i=1}^k f_i - 1\right)} - 1, 0 \right\}. \quad (11)$$

The cutoff $k = 10$ is conventional, and simulation studies suggest that the estimator is reasonably robust to moderate changes in this value [7]. However, for data with extremely heavy tails (as in linguistic applications), the choice of cutoff may have a larger effect and must be studied.

3.4 Jackknife estimators

The jackknife estimator is a general-purpose technique to reduce the bias of an estimator [3]. The idea is as follows: from the reference sample of n individuals, one systematically removes one individual at a time and recomputes the estimator for each reduced sample. The resulting *pseudovalue*s are then combined to produce a bias-corrected estimate. This is the first-order jackknife; the second-order version extends the procedure by considering pairs of removed individuals, yielding a more aggressive bias correction at the cost of higher variance.

The first-order jackknife estimator for species richness is

$$\hat{S}_{jk1} = S_{\text{obs}} + \frac{n-1}{n} \cdot f_1 \approx S_{\text{obs}} + f_1, \quad (12)$$

and the second-order jackknife estimator is

$$\hat{S}_{jk2} = S_{\text{obs}} + \frac{2n-3}{n} \cdot f_1 - \frac{(n-2)^2}{n(n-1)} \cdot f_2 \approx S_{\text{obs}} + 2f_1 - f_2. \quad (13)$$

The first-order jackknife adds approximately f_1 to S_{obs} , effectively assuming that every singleton represents one missed species; this was initially the idea of Turing. It is a rough correction that tends to underestimate the true richness when the sample is small (because many species are missed entirely and do not even appear as singletons) and to overestimate when the sample is large (because some singletons are genuinely rare species that would not yield additional undetected species). The second-order jackknife partially addresses this by subtracting a term proportional to f_2 , which tempers the correction when doubletons are numerous [3].

3.5 Summary and limitations

Each of the estimators above has characteristic strengths and weaknesses. Good-Turing provides the conceptual foundation (sample coverage and probability mass of unseen species), but does not directly produce an estimate of f_0 as a count. Chao1 gives a guaranteed lower bound that is easy to compute and widely applicable but can severely underestimate the true richness when the abundances of species are highly heterogeneous. Precisely the situation we face

with word frequency data. On the other hand, iChao1 tightens the bound by incorporating f_3 and f_4 , but remains a lower bound. ACE accounts for heterogeneity through the coefficient of variation, but depends on the choice of cutoff k and the jackknife estimators offer bias reduction but lack the theoretical guaranty of a lower bound and can overshoot.

A common limitation shared by all of these methods is that they use only the first few frequency counts and treat the rest of the abundance distribution as a black box. In the next section, we explore a fundamentally different approach: instead of extracting information from a handful of frequency counts, we impose a *shape constraint* on the entire distribution of abundances. This leads to the k -monotone and completely monotone estimators, which can exploit more of the available information while remaining nonparametric.

4 Beyond Frequency-Count Estimators

4.1 Nonparametric mixture models

The estimators presented in Section 3 rely on only the first frequency counts $f_1, f_2, \dots, f_r$ for small r , ignoring the rest of the observed abundance distribution. Mixture models take a different approach: they model the *entire* distribution of species abundances, using all available data to estimate the probability $p(0; G)$ that a species goes undetected. Nonparametric mixture models are particularly attractive because they make no assumption about the shape of the underlying abundance distribution, relying only on the assumption that abundances are drawn from some common distribution.

The mixture framework. We begin by defining the mixture model framework. Suppose that we have N species and each species i has its own abundance parameter λ_i , which governs how frequently it appears in a sample. Since we cannot know all individual parameters N , we assume that $\lambda_1, \lambda_2, \dots, \lambda_N$ are independent draws from a common distribution G , called the *mixing distribution*. The model then has two layers:

1. **Conditional layer:** Given λ_i , the number of times species i are observed in the sample follows some simple distribution, typically a Poisson(λ_i) distribution.
2. **Mixing layer:** The parameters $\lambda_1, \dots, \lambda_N$ are themselves random, drawn independently of G .

The mixture distribution. Since each species has its own rate λ_i , we never observe these rates directly (only counts). To obtain the *mixture distribution* (the probability that a randomly chosen species is observed exactly x times in the sample), we need to average the Poisson probability over all possible values of λ , weighted by the probability that each value is under G . Concretely, the Poisson probability of observing x occurrences given rate λ is $\lambda^x e^{-\lambda}/x!$. And multiplying this by $dG(\lambda)$ and integrating over all $\lambda > 0$ gives the unconditional probability:

$$p(x; G) = \int_0^\infty \frac{\lambda^x e^{-\lambda}}{x!} dG(\lambda), \quad x \in \mathbb{N}_0. \quad (14)$$

This is no longer a Poisson distribution; it is a weighted average of Poisson distributions with different rates, where the weights come from G . When we set $x = 0$, it gives us the probability that a species is completely missed in the sample:

$$p(0; G) = \int_0^\infty e^{-\lambda} dG(\lambda). \quad (15)$$

Once $p(0; G)$ is known, the total number of species can be estimated directly.

The mixing distribution G captures the heterogeneity among species abundances. Its shape determines how many species are rare versus common, which in turn determines f_0 , the number of species missed by the sample. For spoken vocabulary data, this heterogeneity is extreme: a handful of function words dominate the token counts.

Why Poisson sampling. The preceding section took the Poisson conditional model as given. We now justify this choice. The Poisson assumption arises naturally from the sampling setup. When recording a speaker for a fixed period of time, each specific word has a small probability of appearing at any given moment, but there are many opportunities for it to occur. This regime, where we have many independent opportunities combined with a small probability per opportunity, is precisely where the Poisson distribution applies. More precisely, the underlying model is binomial, and the Poisson arises as its limit when the number of opportunities grows large. We can derive this connection in two complementary settings: *incidence data*, where the observation period is cut into time slots, and *abundance data*, where time is treated as continuous.

Incidence data: Suppose that we divide a recording of length T into n time slots. In each slot, a given word appears or does not appear. This means that each slot is approximately an independent Bernoulli trial. Let p denote the probability that the word fills any given slot and define $\lambda = np$ as the expected number of times the word is heard in the whole recording. The count X of how many times the word appears then follows a binomial distribution:

$$X \sim \text{Binomial}(n, p), \quad \Pr(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}.$$

Substituting $p = \lambda/n$, the probability mass function becomes

$$\Pr(X = k) = \binom{n}{k} \left(\frac{\lambda}{n}\right)^k \left(1 - \frac{\lambda}{n}\right)^{n-k}.$$

We now take $n \rightarrow \infty$ with λ and k fixed. The combinatorial factor satisfies

$$\binom{n}{k} \left(\frac{\lambda}{n}\right)^k = \frac{n(n-1) \cdots (n-k+1)}{k!} \frac{\lambda^k}{n^k} = \frac{\lambda^k}{k!} (1) \left(1 - \frac{1}{n}\right) \cdots \left(1 - \frac{k-1}{n}\right) \rightarrow \frac{\lambda^k}{k!}, \text{ as } n \rightarrow \infty.$$

For the remaining factor, the classical exponential limit gives

$$\left(1 - \frac{\lambda}{n}\right)^{n-k} = \left(1 - \frac{\lambda}{n}\right)^n \left(1 - \frac{\lambda}{n}\right)^{-k} \rightarrow e^{-\lambda} \cdot 1 = e^{-\lambda}, \text{ as } n \rightarrow \infty$$

since $n - k \approx n$ for large n . Combining the two limits yields the following result.

$$\Pr(X = k) \rightarrow \frac{\lambda^k e^{-\lambda}}{k!},$$

which is the Poisson probability mass function with parameter λ . Thus, for large n , the binomial model is well approximated by a Poisson distribution.

abundance data: The same result can be obtained in a continuous-time framework. If we model the occurrences of each word as a Poisson process with rate λ_i (a counting process satisfying that the rate parameter λ is constant over time, it is independent over non-overlapping intervals, and it is impossible to have simultaneous events), then the number of occurrences in an observation window of length T follows a $\text{Poisson}(\lambda_i T)$ distribution. The derivation proceeds through a system of ordinary differential equations solved by induction; see, e.g., Joyce [14]. Since both routes lead to the same Poisson model, the choice between them is a modeling convenience: the discrete-slot view is more natural for a fixed-length corpus, while the continuous-time view arises naturally when the sampling effort is measured in time rather than token count.

Parametric versus nonparametric mixtures. If G is assumed to belong to a known parametric family, the estimation problem is reduced to fitting a small number of parameters. This can be efficient when the assumed family is correct, but different parametric models can fit the observed data equally well and produce different estimates of S [5]. To illustrate this, we fit three standard two-parameter mixture models to the Shakespeare corpus.

- the *Gamma–Poisson* (negative binomial) mixture, where $\Lambda \sim \text{Gamma}(r, r/\mu)$;
- the *LogNormal–Poisson* mixture, where $\log \Lambda \sim \mathcal{N}(\mu, \sigma^2)$;
- the *Poisson–Inverse Gaussian* mixture, where $\Lambda \sim \text{InvGaussian}(\mu, \phi)$.

In each case, $X \mid \Lambda = \lambda \sim \text{Poisson}(\lambda)$. The full probability mass functions are given in Appendix A.

Figure 3 compares the three fits. The left panel shows the observed region ($x \geq 1$): all three models capture the shape of the abundance distribution reasonably well, with the Poisson–Inverse Gaussian mixture providing the closest fit ($\chi^2 = 79.5$) and the Gamma–Poisson the worst fit ($\chi^2 = 8952.3$). However, the right panel reveals dramatic disagreement in the unobserved region. The Gamma–Poisson model assigns $p(0) = 0.974$, implying that nearly all species are undetected and yields an estimate of $\hat{N} \approx 1,004,540$, roughly 40 times the observed vocabulary. The LogNormal–Poisson gives $\hat{N} \approx 36,788$, while the Poisson–Inverse Gaussian gives $\hat{N} \approx 57,903$. Three models, a comparable fit to the observed data, and estimates of total vocabulary size spanning two orders of magnitude.

This sensitivity to the parametric form of G motivates the nonparametric approach: if G is left completely unspecified, no distributional form is assumed, avoiding the risk of misspecification. This flexibility makes nonparametric mixtures particularly suited to problems like vocabulary estimation, where the true shape of the abundance distribution is unknown.

To estimate the total number of species under this nonparametric framework, we need a likelihood function that relates the observed frequency counts to the unknowns N and G . We derive this by treating the abundances of all N species




Figure 3: Comparison of three parametric Poisson mixture models fitted to the Shakespeare corpus (25,700 observed types, 890,426 tokens). Left: fitted versus observed proportions for $x \geq 1$. Right: full mixture PMF including the extrapolated probability at $x = 0$ (highlighted). The goodness-of-fit summary is shown below the plots.

(observed and unobserved) as independent draws from the mixture distribution $p(x; G)$.

The likelihood. In the mixture model, the abundances of the N species are independent draws from a distribution with probability mass function $p(x; G)$. Since $n_0 = N - S_{\text{obs}}$ species are unobserved, the full likelihood function for N and G based on all species (observed and undetected) is [8]

$$L(N, G) = \frac{N!}{(N - S_{\text{obs}})! \prod_{x \geq 1} f_x!} \{p(0; G)\}^{N - S_{\text{obs}}} \prod_{x \geq 1} \{p(x; G)\}^{f_x}. \quad (16)$$

The first factor $\{p(0; G)\}^{N - S_{\text{obs}}}$ accounts for the $N - S_{\text{obs}}$ species that were never observed, each independently having a probability $p(0; G)$ of being missed. The second factor accounts for the observed frequency counts. Both N and G are unknown and must be estimated jointly.

Sanathanan [16] showed that the full likelihood decomposes as $L(N, G) = L_m(N, G) \times L_c(G)$, where L_m is a binomial likelihood for S_{obs} given N and $p(0; G)$, and L_c is the conditional likelihood of the observed frequency counts given S_{obs} , based on the distribution of the zero-truncated mixture. This decomposition gives rise to two estimation approaches:

- The **conditional approach** first estimates G by maximizing $L_c(G)$ (fit-

ting the zero-truncated data), then estimates N from L_m via

$$\hat{N} = \left\lfloor \frac{S_{\text{obs}}}{1 - p(0; \hat{G})} \right\rfloor. \quad (17)$$

- The **unconditional approach** jointly maximizes $L(N, G)$ in both N and G . Wang and Lindsay [20] showed that the resulting estimate also satisfies Equation (17).

In both cases, the estimation of N reduces to the estimation of $p(0; G)$ well. This makes the choice of mixture model critical, particularly its behavior near $\lambda = 0$.

Likelihood decomposition. Each of the N species is independently observed (with probability $1 - p(0; G)$) or missed (with probability $p(0; G)$). The number of observed species therefore follows a binomial distribution.

$$S_{\text{obs}} \sim \text{Binomial}(N, 1 - p(0; G)).$$

This lets us factor in the full likelihood (16) as $L(N, G) = L_m(N, G) \times L_c(G)$ [16], where

$$L_m(N, G) = \binom{N}{S_{\text{obs}}} [1 - p(0; G)]^{S_{\text{obs}}} [p(0; G)]^{N - S_{\text{obs}}} \quad (18)$$

is the marginal likelihood of observing exactly S_{obs} species, and

$$L_c(G) = \frac{S_{\text{obs}}!}{\prod_{x \geq 1} f_x!} \prod_{x \geq 1} \left[\frac{p(x; G)}{1 - p(0; G)} \right]^{f_x} \quad (19)$$

is the conditional likelihood of the observed frequency counts given S_{obs} . The cell probabilities in L_c are precisely those of the zero-truncated mixture distribution:

$$p^+(x; G) = \frac{p(x; G)}{1 - p(0; G)}, \quad x \geq 1, \quad (20)$$

so that, conditional on S_{obs} , the observed abundances behave as S_{obs} independent draws from $p^+(\cdot; G)$. The decomposition can be verified by multiplying (18) and (19) and recovering (16).

The conditional approach and the estimator $\hat{N}$. The conditional approach proceeds in two stages. First, estimate the mixing distribution by maximizing $L_c(G)$ over G , producing $\hat{G}_c$. Since L_c depends only on zero-truncated probabilities, this step uses only the observed frequency counts and does not involve N . Second, with $\hat{G}_c$ fixed, estimate N from marginal likelihood $L_m(N, \hat{G}_c)$. Writing $q = 1 - p(0; \hat{G}_c)$ for the estimated detection probability, L_m becomes a binomial likelihood in a single unknown N :

$$L_m(N) = \binom{N}{S_{\text{obs}}} q^{S_{\text{obs}}} (1 - q)^{N - S_{\text{obs}}}.$$

To find the integer maximizer, consider the likelihood ratio of consecutive values:

$$\frac{L_m(N+1)}{L_m(N)} = \frac{\binom{N+1}{S_{\text{obs}}}}{\binom{N}{S_{\text{obs}}}} \cdot (1-q) = \frac{N+1}{N+1-S_{\text{obs}}} \cdot (1-q).$$

This ratio exceeds 1 (i.e., the likelihood is still increasing) if and only if

$$(N+1)(1-q) > N+1-S_{\text{obs}} \iff N < \frac{S_{\text{obs}}}{q} - 1.$$

Therefore, the likelihood increases for all $N < S_{\text{obs}}/q - 1$ and decreases once $N \geq S_{\text{obs}}/q - 1$, giving the integer-valued maximum likelihood estimator

$$\hat{N} = \left\lfloor \frac{S_{\text{obs}}}{q} \right\rfloor = \left\lfloor \frac{S_{\text{obs}}}{1-p(0; \hat{G}_c)} \right\rfloor. \quad (21)$$

The floor function arises because N must be an integer: when S_{obs}/q is not an integer, the likelihood peaks between $\lfloor S_{\text{obs}}/q \rfloor$ and $\lceil S_{\text{obs}}/q \rceil$, and the ratio test shows that the likelihood is maximized at the largest integer below S_{obs}/q .

This estimator has a natural interpretation as a Horvitz–Thompson-type correction: we observe S_{obs} species, each detected with probability q , so the total population is approximately S_{obs}/q . The quality of $\hat{N}$ depends entirely on how well $p(0; \hat{G}_c)$ is estimated, which in turn depends on the choice of the mixture model and its behavior near $\lambda = 0$.

The boundary problem. The difficulty is that $p(0; G) = \int_0^\infty e^{-\lambda} dG(\lambda)$ is sensitive to the masses G near $\lambda = 0$: species with very small rates contribute the most to $p(0; G)$, but they are precisely the species least likely to appear in the observed data. Since zero-truncated frequencies barely constrain G in this region, small differences in the estimated shape of G near zero can shift $p(0; \hat{G})$ substantially and because $\hat{N} = S_{\text{obs}}/(1-p(0; \hat{G}))$, even a modest increase in $p(0; \hat{G})$ toward 1 causes $\hat{N}$ to explode.

4.2 Frequency-ratio regression (Breakaway)

The Breakaway method, developed by Willis and Bunge [21], takes a different approach to estimating species richness: instead of fitting a distributional model to the abundances, it works directly with the ratios of consecutive frequency counts. Crucially, it does not assume a mixed Poisson model, departing from the framework of Section 4.1.

The classical estimators presented in Section 3 rely only on f_1 and f_2 (or up to f_4 for iChao1), and can perform poorly on datasets characterized by a large number of rarely observed species, a high singleton count, and a long tail of very abundant species, precisely the pattern exhibited by word frequency data. Breakaway addresses this by using many frequency ratios $f_2/f_1, \dots, f_{j+1}/f_j, \dots$ to fit a smooth curve, then extrapolating back to $j = 0$ to estimate f_0 . This borrows strength from the entire tail of the distribution rather than relying on the first few counts alone.

How the method works. Let $f_1, f_2, \dots, f_n$ be the observed frequency counts. Since the total number of species satisfies $S = f_0 + \sum_{j=1}^n f_j = f_0 + S_{\text{obs}}$, the problem is reduced to estimating as previously f_0 . Breakaway fits a rational function of the form

$$\frac{f_{j+1}}{f_j} = \frac{\beta_0 + \beta_1 j}{1 + \beta_2 j + \beta_3 j^2} \quad (22)$$

to the observed ratios using heteroscedastic nonlinear regression.¹ Setting $j = 0$ in the fitted model gives the predicted ratio $f_1/\hat{f}_0 = \hat{\beta}_0$, so that

$$\hat{f}_0 = \frac{f_1}{\hat{\beta}_0}, \quad \hat{S} = S_{\text{obs}} + \hat{f}_0. \quad (23)$$

In practice, Breakaway does not fit a single model, it considers a family of rational functions of increasing complexity and selects the most simple one satisfying three conditions: $\hat{\beta}_0 > 0$ (so that $\hat{f}_0$ is positive), no singularities in the domain of the observed ratios, and numerical convergence [21]. The issue of unobserved negative diversity estimators is as old as the problem itself: Fisher, Corbet, and Williams (1943) recognized it by fitting the negative binomial distribution to the Malayan butterfly data [10].

Specifically, since $\hat{f}_0 = f_1/\hat{\beta}_0$, the approximate variance of $\hat{f}_0$ is

$$\widehat{\text{Var}}(\hat{f}_0) \approx f_1 \hat{\beta}_0^{-2} \left(1 - \frac{f_1}{n} + f_1 \hat{\beta}_0^{-2} \widehat{\text{Var}}(\hat{\beta}_0) \right), \quad (24)$$

where $\widehat{\text{Var}}(\hat{\beta}_0)$ is the estimated variance of $\hat{\beta}_0$ from nonlinear regression, and the variance of the total estimate is [21]

$$\widehat{\text{Var}}(\hat{S}) \approx n\hat{f}_0/(n + \hat{f}_0) + \widehat{\text{Var}}(\hat{f}_0) \quad (25)$$

Figure 4 illustrates the extrapolation step in the Breakaway method. The observed points are the consecutive frequency ratios f_{j+1}/f_j , calculated from the abundance frequency counts. We can see the fit of the smooth rational function and the fit value in the estimates of $j = 0$ $\hat{\beta}_0 = f_1/\hat{f}_0$, so the number of unseen types is obtained as $\hat{f}_0 = f_1/\hat{\beta}_0$. The different curves show how the fit changes when more ratio terms are included.

Limitations. As we have seen previously, we have no guaranty that the estimate is positive. If $\hat{\beta}_0$ is close to 0 or negative, The estimate $\hat{f}_0 = f_1/\hat{\beta}_0$ will have a negative or very high estimate.

This method also requires enough data points; otherwise we will not have the desired frequency counts. As stated in the Willis paper [21], if the maximum number of observed frequency is 6 or less, then there aren't enough ratios to fit the nonlinear model, and Breakaway cannot run.

¹That is, the regression accounts for the fact that the variance of the errors is not constant across the data points.




Figure 4: Breakaway frequency-ratio regression for one example corpus item. Black points show the observed ratios f_{j+1}/f_j . The colored curves are rational-function fits using ratios up to different maximum indices J . Their extrapolated values at $j = 0$ determine $\hat{\beta}_0$, and hence the estimate $\hat{f}_0 = f_1/\hat{\beta}_0$.

4.3 Shape-constrained estimation

4.3.1 k -monotone distributions: definition and mixture representation

The nonparametric Poisson mixture model, as we have seen, reduces species richness estimation to the problem of estimating $p(0; G)$. However, it suffers from the boundary problem: the NPMLE (nonparametric Maximum Likelihood Estimator) of G tends to place mass near $\lambda = 0$, which inflates $p(0; G)$ and sends $\hat{N}$ to very high values. This motivates the move to an alternative nonparametric mixture model that retains the flexibility of leaving the mixing distribution unspecified, but does not suffer from this instability.

The following strategy follows those ideas: instead of constraining the mixing distribution G , we impose a shape constraint directly on the abundance PMF p . We require that p belongs to the family of k -monotone distributions:

The definition of k -monotone distributions is given as follows: Let k be a positive integer. The *forward difference operator* Δ acts on a sequence $p: \mathbb{N}_0 \rightarrow \mathbb{R}$ by

$$\Delta p(x) = p(x+1) - p(x),$$

higher-order differences are defined recursively: the a th order forward difference

is

$$\Delta^a p(x) = \sum_{i=0}^a \binom{a}{i} (-1)^i p(x+a-i), \quad a = 1, 2, \dots$$

A discrete distribution with probability mass function p in $\mathbb{N}_0 = \{0, 1, 2, \dots\}$ is said to be *k-monotone* if

$$(-1)^a \Delta^a p(x) \geq 0 \quad \text{for all } a = 1, \dots, k \text{ and all } x \in \mathbb{N}_0.$$

If the condition holds for all $a \in \mathbb{N}$ (i.e., $k = \infty$), the distribution is called *completely monotone*.

For $k = 1$, the condition reduces to $-\Delta p(x) = p(x) - p(x+1) \geq 0$ for all x , i.e. p does not increase. For $k = 2$, it additionally requires $\Delta^2 p(x) = p(x+2) - 2p(x+1) + p(x) \geq 0$, meaning that p is non-increasing *and* convex: the successive decrements $p(x) - p(x+1)$ are themselves non-increasing. Both conditions are natural for species abundance distributions, which are typically right-skewed with their mode at $x = 1$ and a long decreasing tail.

The key structural property of k -monotone distributions is that they admit a mixture representation. Lefevre and Loisel [15] showed that a PMF p in $\mathbb{N}_0$ is k -monotone if and only if there exists a sequence of mixing probabilities $(\pi_l^k)_{l \in \mathbb{N}_0}$ with $\pi_l^k \geq 0$ and $\sum_{l \geq 0} \pi_l^k = 1$ such that

$$p(j) = \sum_{l \geq 0} \pi_l^k Q_l^k(j),$$

where

$$Q_l^k(j) = \frac{\binom{l-j+k-1}{k-1}}{\binom{l+k}{l}} \mathbf{1}_{\{0 \leq j \leq l\}}$$

is the *k-monotone spline* with support $\{0, 1, \dots, l\}$. Each spline Q_l^k is itself a discrete distribution concentrated in $l+1$ points, and the mixing weights π_l^k describe how the overall abundance distribution is assembled from these components. The weights can be recovered from p using the inversion formula

$$\pi_l^k = (-1)^k \Delta^k p(l) \binom{l+k}{l}.$$

This mixture representation is the analog of the Poisson mixture in Section 4.1, but with the k -monotone splines Q_l^k replacing the Poisson distributions as building blocks. The crucial difference is that the spline Q_0^k (with $l = 0$) is a point mass at the origin: it assigns all its probability to $j = 0$. In the context of species richness, a species whose abundance follows Q_0^k would never be observed in any sample. Since something that was never observed cannot contribute to the observed data, we must set $\pi_0^k = 0$. Applying the inversion formula with $l = 0$ gives $\pi_0^k = (-1)^k \Delta^k p(0) = 0$, which is equivalent to the constraint

$$\Delta^k p(0) = p(0) + \sum_{j=1}^k (-1)^j \binom{k}{j} p(j) = 0.$$

This single linear equation links $p(0)$ (the probability that a species goes completely undetected) to the observed probabilities $p(1), p(2), \dots, p(k)$. Rearranging in terms of the zero-truncated PMF $p^+(j) = p(j)/(1 - p(0))$ yields

$$\frac{1}{1 - p(0)} = 1 - \sum_{j=1}^k (-1)^j \binom{k}{j} p^+(j).$$

This is precisely what the Poisson mixture framework lacked: a structural constraint that identifies $p(0)$ from the observed data. The boundary problem is resolved not by penalizing the mass near $\lambda = 0$, but by working with a model in which $p(0)$ is determined by the shape of the observed distribution. Once $p(0)$ is identified, the estimator $\hat{N} = \lfloor S_{\text{obs}}/(1 - p(0)) \rfloor$ follows directly. This solves the problem of *one cannot see what cannot be seen*.

Figure 5 gives a visual intuition of the role of the parameter k . We see that large values of k impose alternating-sign constraints, thus increasing k restricts the possible shapes of the abundance distribution, making the curve more structured.




Figure 5: Illustrative probability mass functions satisfying different levels of k -monotonicity. The $k = 1$ curve is a decreasing PMF, the $k = 2$ and $k = 4$ curves are k -monotone splines, and the $k = \infty$ curve is a geometric PMF, which is completely monotone.

4.3.2 Selection of k

The value of k controls the shape restriction, a larger k imposes more constraints and therefore structure and pushes the support points of the NPMLE further from the left boundary. It is important to note that it does not necessarily produce a higher fitted probability at 0. Therefore, the estimate of $\hat{N}$ does not necessarily increase with k [8].

In practice, the appropriate value of k for a given data set is unknown, resulting in a model selection problem. Since the family of k -monotone distributions shrinks as k increases, the maximized log-likelihood $\ell_k(\hat{N}, \hat{G})$ decreases monotonically with k . Chee and Wang balance the goodness of fit against model complexity by evaluating the Akaike information criterion (AIC) over a grid of k values ranging from $k = 2$ to $k = \infty$, selecting the value that minimizes AIC. The Akaike information criterion is defined as

$$\text{AIC}_k = -2 \ell_k(\hat{N}, \hat{G}) + 2 d_k, \quad (26)$$

where d_k denotes the number of support points of the NPMLE $\hat{G}$ under k -monotonicity, which serves as the effective number of free parameters in the mixing distribution.

4.3.3 Complete monotonicity as the limit of k -monotonicity

As k increases, the family of k -monotone distributions becomes increasingly restrictive. A classical result given by Steutel [19] characterizes the limiting case: a PMF p on $\mathbb{N}_0$ is completely monotone if and only if it is a mixture of geometric distributions, that is,

$$p(j) = \int_0^1 \alpha(1 - \alpha)^j dH(\alpha), \quad j \in \mathbb{N}_0 \quad (27)$$

where H is a distribution function on $[0, 1]$. Thus, geometric distributions play the same role for complete monotonicity as the k -monotone splines Q_t^k play for finite k , i.e. they are the building blocks from which any completely monotone PMF can be assembled.

Chee and Wang [8] explicitly established the connection by showing that the k -monotone mixture PMF converges, as in $k \rightarrow \infty$, to a mixture of geometric distributions. Balabdaoui and Kulagina [2] formalized this further: if a sequence of k -monotone PMF $(p^k)_{k \in \mathbb{N}}$ has mixing distributions that converge weakly to some distribution H on $[0, 1]$, then the sequence converges pointwise to the completely monotone PMF (27).

With this convergence result, we can justify using the completely monotone model whenever the true degree of monotonicity is large. In practice, when the AIC selection of Chee and Wang keeps selecting larger values of k , this suggests that the data is approximately completely monotone, and one can work directly with the limiting model. Doing so eliminates the need to select k altogether, which is a significant practical advantage, as the differences between k and $(k + 1)$ -monotonicity become increasingly difficult to detect as k grows.

4.3.4 The Balabdaoui and Kulagina estimator: $\hat{N}_\infty$

Balabdaoui and Kulagina [2] avoid the problem of selecting k by letting it grow logarithmically with the sample size n . Define $\bar{\alpha}_n = 1 - \bar{p}_n(1)$, where $\bar{p}_n$ is the empirical estimator of the zero-truncated PMF. The degree of monotonicity is set to

$$k_n = \left\lfloor \frac{\log n}{2 \log(\bar{\alpha}_n^{1/2} + 1)} \right\rfloor,$$

and the estimator of N is

$$\hat{N}_\infty = \left\lfloor n \left(1 - \sum_{j=1}^{k_n} (-1)^j \binom{k_n}{j} \bar{p}_n(j) \right) \right\rfloor.$$

Under the assumption that $\bar{\alpha}_n \geq (\sqrt{2} - 1)^2$ and that the true PMF is a finite mixture of geometrics, the estimator is asymptotically normal, resulting in the confidence interval at level $1 - \alpha$:

$$\hat{N}_\infty \pm z_{1-\alpha/2} \sqrt{n} \hat{\sigma}_n, \quad \hat{\sigma}_n^2 = \sum_{l=1}^{k_n} \binom{k_n}{l}^2 \bar{p}_n(l).$$

4.4 Summary

Section 4 has developed a progression of increasingly structured approaches to estimate the count of unobserved species f_0 . The classical Poisson mixture framework of Section 4.1 reduces the problem to estimating $p(0; G)$, but suffers from the boundary instability that sends $\hat{N}$ to unreliable values. Breakaway (Section 4.2) avoids this by working directly with frequency ratios, at the cost of numerical instability for small samples. Shape-constrained estimation (Section 4) resolves the boundary problem structurally: by imposing k -monotonicity on the abundance PMF, $p(0)$ is identified directly from the shape of the observed data, and the Balabdaoui–Kulagina estimator $\hat{N}_\infty$ eliminates the need to select k altogether. With all estimators now defined, we turn to the data. Section 5 describes the corpora and preprocessing decisions, and Section 6 applies all methods to spoken language data in English, French, and German, as well as to the Shakespeare corpus as a written replication baseline.

5 Data and Methodology

5.1 What is a word?

Before implementing our models, we must first specify what we mean by a *word*. This is not entirely straightforward, since words can be counted in different ways depending on the linguistic conventions adopted. The issue becomes particularly important when comparing texts across different languages, as languages differ substantially in how words are formed. For example, German productively forms compound words, allowing the creation of long expressions that may be considered distinct lexical items.

The first distinction is between *tokens* and *types*. A token refers to each individual occurrence of a word in a text, whereas a type refers to a distinct word form. For example, in the sentence

the cat and the dog,

there are five tokens, but only four types, since the word *the* appears twice.

Another important distinction concerns whether different grammatical forms of a word should be treated as separate types. For example, one may choose either:

- treat *run*, *running*, and *ran* as different words, or
- group them together under the same base form.

The latter approach is known as *lemmatization*. A *lemma* is the canonical or dictionary form of a word, so conjugated or plural forms are counted together. Under this convention, *run*, *running*, and *ran* are all associated with the lemma *run*.

In this work, we consider the approach of counting tokens. This method was used for estimating the well known example about vocabulary richness of *Shakespeare*. And because of convenience and for the sake of comparison, we will make use of the same.

5.2 Data sources

Our analysis draws on five data sources: an established reference corpus, three sources of spoken or dialog-based English, one of spoken French, and one of spoken German.

Shakespeare corpus. As a baseline for replication and comparison, we use the word frequency counts from Spevack’s concordance [18], comprising 890,426 tokens and 25,700 observed types across the complete works of Shakespeare. This data set has been widely used in the literature on species richness [2, 8] and allows us to verify our implementations against published results.

Internet Movie Script Database (IMSDb). We downloaded scripts from approximately 1,200 English-language films from the Internet Movie Script Database. We applied preprocessing tools to extract only the dialog portions, discarding stage directions and scene descriptions. The lines of each character were then aggregated throughout the film. To ensure a minimum sample size, we retained only characters with at least 500 tokens, corresponding to roughly 3–5 minutes of speech at a comfortable speaking pace of 100–150 words per minute.

British National Corpus (BNC). We used the spoken component of the British National Corpus, which contains transcriptions of spontaneous British English conversation and other spoken registers.

Santa Barbara Corpus of Spoken American English. This corpus consists of transcriptions of naturally occurring spoken interactions in American English.

CLAIRE Dialogue Dataset. We used the CLAIRE dataset, a large collection of French conversational data aggregated from various dialog sources.

Datenbank Gesprochenes Deutsch (DGD). We used the DGD corpus, a database of spoken German transcribed maintained by the Leibniz Institute for the German Language (IDS Mannheim).

Our analysis combines one reference literary corpus with five corpora of spoken or dialog-based language.²

5.3 Preprocessing

All texts were lowercased and devoid of punctuation. We remove all the numbers and filler words and leave the contractions as a single type. For each speaker or character, we computed the abundance frequency counts $f_1, f_2, \dots$ as defined in Section 2: f_k is the number of types that appear exactly k times in the text of that speaker. These frequency counts serve as input for all estimators described in Sections 3 and 4.

²See Appendix B for URLs and access information for all data sources.

6 Results

This section presents the results in four stages. We begin with the Shakespeare corpus as a replication baseline to validate the estimator implementations, then characterize the spoken corpora by language, and compare all estimators within and across languages. Finally, we examine the shape-constrained estimates. Throughout, we use normalized quantities such as $\hat{S}/S_{obs}$ and sample coverage to account for the substantial differences in corpus size between languages.

6.1 Shakespeare replication baseline

We first use the full Shakespeare corpus as a replication baseline for the estimators. This corpus is not used in the cross-language spoken language comparison, since it is a written literary corpus rather than spoken language.

x	1	2	3	4	5	6	7	8	9	10
0+	9704	3562	2050	1408	1028	759	620	475	415	393
10+	292	264	217	231	192	176	165	143	144	146
20+	124	103	89	97	87	89	76	84	59	82
30+	74	49	60	56	57	51	34	36	42	40
40+	43	52	33	34	23	36	33	33	35	34
50+	24	17	20	34	27	22	13	21	20	18
60+	25	17	18	23	23	13	13	10	16	13
70+	16	8	14	17	9	17	8	14	17	11
80+	14	16	10	15	8	14	9	6	11	10
90+	3	9	8	12	8	5	8	10	6	5

Table 1: Frequency spectrum of the processed full Shakespeare corpus. Each entry is the number of distinct word types occurring exactly x times, where the row label gives the tens digit; for example, the first row gives $f_1, \dots, f_{10}$ and the second row gives $f_{11}, \dots, f_{20}$. The full corpus has 890,426 tokens and 25,700 observed word types.

The frequency spectrum in table 1 shows the typical long-tailed structure of the word-frequency data. A large number of word types occur only once or twice, whereas progressively fewer types occur at higher frequencies. This is important because most richness estimators rely heavily on low-frequency counts, especially f_1 and f_2 .

We observe (Table 2) that the classical lower-bound estimators give in total between 35,000 and 41,500 types, while ACE gives a substantially larger estimate of about 70,900 types. This difference reflects the fact that ACE includes an explicit heterogeneity correction for the rare part of the frequency spectrum. Breakaway is not reported for the full Shakespeare corpus because the fitted extrapolation parameter was negative, making the estimated number of unseen types undefined.

Estimator	Tokens n	S_{obs}	f_1	f_2	Coverage	$\hat{S}$	$\hat{S}/S_{\text{obs}}$
Chao1	890426	25700	9704	3562	0.9891	38918.3	1.5143
iChao1	890426	25700	9704	3562	0.9891	41506.7	1.6150
ACE	890426	25700	9704	3562	0.9891	70942.2	2.7604
Jackknife1	890426	25700	9704	3562	0.9891	35404.0	1.3776
Jackknife2	890426	25700	9704	3562	0.9891	41546.0	1.6166
Breakaway	890426	25700	9704	3562	0.9891	—	—

Table 2: Summary of Shakespeare replication estimates for the processed full Shakespeare corpus. Breakaway is undefined here because the fitted extrapolation parameter is negative.




Figure 6: Rarefaction and extrapolation curve for a representative Shakespeare unit.

Figure 6 shows that the rarefaction and extrapolation curve rises steeply at small sample sizes and then begins to flatten, following a roughly logarithmic shape, but has not fully saturated at the observed sample size. Additional sampling would therefore still be expected to reveal new word types, although at a diminishing rate. The extrapolated portion should be read as a model-based projection rather than as an observed quantity.

6.2 Corpus and language overview

We next summarize the spoken and dialog-based corpora used for the main multilingual analysis. In the cross-language comparison, English consists of IMSDb, BNC, and the Santa Barbara Corpus; French consists of CLAPI; and German consists of DGD.

Language	Corpora included	Number of units	Total tokens	Median tokens	Mean coverage
English	bnc, imsd, sbcorpus	5517	6675363	815	0.7197
French	clapi	10813	160251286	1927	0.8104
German	dgd	47	68141	1485	0.7782

Table 3: Language-level overview of the spoken corpora used in the cross-language comparison. Shakespeare is excluded because it is a written literary corpus.

Table 3 shows that there is an imbalance in the number of tokens per language. French has the highest, followed by English and then German. Even though English has a higher number of tokens, its mean coverage is smaller than that of German by approximately 6%. Therefore, later comparisons should use normalized quantities, such as $\hat{S}/S_{\text{obs}}$ and coverage, rather than relying only on estimated raw richness.

Table 4 exhibits that at the corpus level, CLAPI dominates the total amount of data, while IMSDb contributes most of the English units. The table also shows that some corpora contain very small units, which can make the richness estimates unstable. Shakespeare is included only for context and should not be interpreted as part of the spoken-language comparison.

Figure 7 shows that French has the highest sample coverage across all three languages, indicating that its observed sample is the most complete relative to the estimated total vocabulary. Consequently, French also has the lowest proportion of word types appearing only once, since high coverage is directly linked to a smaller singleton fraction.

The singleton proportions in table 5 reinforce the coverage picture. French has the smallest singleton share relative to the observed vocabulary ($f_1/S_{\text{obs}} = 0.34$), while German has the largest share (0.56). A higher singleton fraction means that a larger share of the next sampled tokens is expected to belong to previously unseen types, which implies lower coverage and stronger required extrapolation. This explains why German requires the most extrapolation despite having the smallest total corpus.

Corpus	Language	Units	Total tokens	Median tokens	Min tokens	Max tokens	Median S_{obs}	Median coverage
BNC	English	136	312275	692	2	36869	282	0.7947
CLAPI	French	10813	160251286	1927	347	2757857	585	0.8210
DGD	German	47	68141	1485	423	2392	543	0.7838
IMSDb	English	5335	6307786	814	61	9529	357	0.7185
SBCorpus	English	46	55302	1156	1	5311	404	0.7930
Shakespeare	Shakespeare	43	1780852	21547	362	890426	3561	0.9094

Table 4: Corpus-level overview of all data sources. Shakespeare is included for context as the written replication baseline and is not part of the spoken-language comparison.

Language	Total tokens	Observed types	f_1	f_2	f_1/S_{obs}	f_1/n
English	6675363	81575	31360	10827	0.3844	0.0047
French	160251286	208726	71965	23787	0.3448	0.0004
German	68141	8024	4453	1334	0.5550	0.0654

Table 5: Low-frequency vocabulary summary by language for the spoken corpora. Here f_1 is the number of word types appearing exactly once and f_2 the number appearing exactly twice. The full comparable frequency spectrum is stored in `results/tables/language_frequency_spectrum_comparable.csv`.




Figure 7: Sample coverage by language for the spoken corpora.




Figure 8: Distribution of token counts by language for the spoken corpora. Each unit corresponds to one speaker file in the processed database. The log scale is used because token counts span several orders of magnitude.

In Figure 8 we can observe that the left panel uses a density y-axis so that French, which contributes far more units than the other languages, does not dominate the display; this allows the three languages to be compared on equal footing. The right panel shows token counts per unit by corpus, where CLAPI has the highest median by a large margin. Together, the two panels confirm that sample size varies substantially across languages and corpora and is therefore a major confounder for all later richness comparisons.

6.3 Per-language estimator results

This section compares the richness estimates produced by the different estimators within each language.

Breakaway fails quite often because it does not have enough consecutive nonzero frequency counts for a stable ratio-regression fit; when the fitted parameter is negative, it is undefined. This is shown in table 6. When comparing normalized estimates between estimators, the ratios are broadly similar across languages, as also visible in Figure 9. Across all three languages, Jackknife1 is the most conservative estimator, Chao1 and Jackknife2 are intermediate, ACE is more aggressive, and Breakaway has the largest median ratio among the successful fits. The normalized ratios are fairly similar across languages: for example, Chao1 gives median ratios around 2.1–2.35, ACE around 3.8–4.1, and Breakaway around 4.9–5.3. This suggests that the estimator choice affects the result more strongly than language alone.

6.4 Completely monotone and k-monotone model selection

The classical estimators in Section 6.3 treat the abundance distribution as a black box, relying on only f_1 and f_2 . We now turn to the shape-constrained es-

Language	Estimator	Mean $\hat{S}$	Median $\hat{S}$	IQR $\hat{S}$	Median $\hat{S}/S_{\text{obs}}$	Failed
English	Chao1	1015.4	885.9	546.6	2.3519	0
English	iChao1	1111.7	968.4	595.0	2.5719	0
English	ACE	1771.5	1523.0	994.7	4.0940	10
English	Jackknife1	694.8	585.6	378.7	1.6352	0
English	Jackknife2	901.2	764.7	486.2	2.1257	2
English	Breakaway	3320.0	1927.2	1880.0	5.2128	2764
French	Chao1	2721.7	1377.2	1968.9	2.1066	0
French	iChao1	2960.7	1507.5	2155.8	2.3069	0
French	ACE	4990.2	2443.5	3514.4	3.8163	0
French	Jackknife1	2155.6	928.8	1400.0	1.5676	0
French	Jackknife2	2677.4	1200.7	1806.1	2.0035	0
French	Breakaway	22309.1	3249.9	7208.9	4.9321	5516
German	Chao1	1163.3	1269.4	406.5	2.2451	0
German	iChao1	1269.6	1359.3	438.1	2.4446	0
German	ACE	2075.6	2188.3	811.5	4.0295	0
German	Jackknife1	808.8	855.8	227.5	1.6027	0
German	Jackknife2	1045.6	1111.3	312.8	2.0642	0
German	Breakaway	5746.8	2500.3	2619.9	5.2887	20

Table 6: Per-language estimator summary for the spoken corpora.




Figure 9: Distribution of estimated richness by estimator and language.

timators developed in Section 4, which impose structure on the full distribution, and ask whether they produce meaningfully different estimates.

The following figures document how the finite- k and completely monotone shape-constrained estimates behave.

Language/Corpus	Units	Median k^*	Mean k^*	$k^* = k_{\max}$ (%)	Median $\hat{S}_{CM}$
English	15	5.0	5.0	0.0	1707.0
French	15	6.0	6.0	0.0	17525.0
German	15	4.0	4.1	0.0	605.0
Shakespeare	43	8.0	7.3	0.0	3390.0

Table 7: Summary of selected finite- k monotone and completely monotone estimates.

In all languages and corpora, the AIC-selected k^* ranges from small to moderate and never reaches the maximum tested value as shown in table 7. This argues that the data does not push toward a completely monotone limit on its own: the finite- k model is adequate and complete monotonicity is not required by the selection criterion. The completely monotone model is therefore useful as a theoretical benchmark, but should not be treated as the uniquely correct choice for these data.

Furthermore the selected values are moderate rather than minimal: $k^* = 5$ for English, $k^* = 6$ for French, $k^* = 4$ for German, and $k^* = 8$ for Shakespeare (Figure 10). This suggests that the data are not well described by very low-order monotonicity, but the optimum also does not occur at the largest tested value. Thus, the figure supports the use of moderately flexible finite- k monotone models, while giving only limited evidence for treating the completely monotone model as a necessary limiting case. The steep decrease in AIC for small values of k shows that increasing the order of monotonicity initially substantially improves the fit. After the selected value, the curve becomes almost flat, so additional complexity gives little improvement and is penalized by AIC.




Figure 10: AIC values across candidate values of k for representative units.

We can see on Figure 11 that the AIC-selected finite- k^* monotone estimates and the completely monotone estimates are in near-perfect agreement: almost all points lie on the diagonal, meaning both approaches produce essentially the same estimated richness. A strong deviation from the diagonal would have




Figure 11: Comparison between completely monotone estimates and selected finite- k monotone estimates.

indicated that the choice between finite- k and complete monotonicity materially affects the estimates, but that is not the case here. In practice, the choice of the model has more impact on the theoretical interpretation than on the numerical result.

6.5 Rarefaction and extrapolation

The estimators in the preceding sections each produce a single point estimate of total richness. Rarefaction and extrapolation curves complement this by showing how the observed vocabulary accumulates as sampling effort increases, and how far each speaker is from saturation. This matters for interpreting the point estimates: a speaker whose curve is still rising steeply at the observed sample size has a much larger unobserved vocabulary than the raw S_{obs} suggests, while a flattening curve indicates that most types have already been seen.




Figure 12: Rarefaction and extrapolation curves for representative spoken units in each language.

All three curves continue to increase at their respective observed sample sizes, confirming that none of the representative units have reached saturation (Figure 12). The curves also converge toward broadly similar extrapolated values, though this should not be taken as evidence that the languages have equal vocabulary richness: each panel represents a single selected unit, and the corpora differ substantially in size and composition.

6.6 Cross-language comparison

This comparison uses only spoken and dialog-based corpora. Shakespeare is excluded here because it is a written literary corpus and serves only as the replication baseline in the first results subsection.

Unlike Table 6, which reported the mean and IQR of $\hat{S}$ together with failure counts, the table 8 focuses on the median observed vocabulary, the median estimated total richness, the median unseen component $\hat{f}_0$ and the normalized ratio $\hat{S}/S_{obs}$. The raw estimates are highest for French, but the normalized ratios are much closer across languages, which shows why the raw $\hat{S}$ values should not be compared without considering the sample size and coverage.

Language	Estimator	Median S_{obs}	Median $\hat{S}$	Median $\hat{f}_0$	Median coverage	Median $\hat{S}/S_{\text{obs}}$
English	Chao1	357.0	885.9	520.1	0.7194	2.3519
English	iChao1	357.0	968.4	605.8	0.7194	2.5719
English	ACE	357.0	1523.0	1161.7	0.7194	4.0940
English	Jackknife1	357.0	585.6	228.2	0.7194	1.6352
English	Jackknife2	357.0	764.7	405.9	0.7194	2.1257
English	Breakaway	357.0	1927.2	1554.4	0.7194	5.2128
French	Chao1	585.0	1377.2	797.9	0.8210	2.1066
French	iChao1	585.0	1507.5	929.2	0.8210	2.3069
French	ACE	585.0	2443.5	1855.0	0.8210	3.8163
French	Jackknife1	585.0	928.8	342.8	0.8210	1.5676
French	Jackknife2	585.0	1200.7	614.5	0.8210	2.0035
French	Breakaway	585.0	3249.9	2611.2	0.8210	4.9321
German	Chao1	543.0	1269.4	739.4	0.7838	2.2451
German	iChao1	543.0	1359.3	831.2	0.7838	2.4446
German	ACE	543.0	2188.3	1645.3	0.7838	4.0295
German	Jackknife1	543.0	855.8	312.8	0.7838	1.6027
German	Jackknife2	543.0	1111.3	568.3	0.7838	2.0642
German	Breakaway	543.0	2500.3	1940.1	0.7838	5.2887

Table 8: Cross-language estimator summary for spoken corpora only.




Figure 13: Median estimated richness by language and estimator for the spoken corpora.

As observed in Figure 13, the cross-language difference that is consistent between estimators is that French has the highest median $\hat{S}$, followed by German and then English. This is not only because French has a larger corpus: the English corpus is larger than the German one, and still scores lower, suggesting that corpus size alone does not drive the ordering.

Also, when comparing the normalized ratios rather than the raw estimates, the French corpus no longer stands out as clearly (Figure 14). This is mainly because the number of estimated word types does not grow linearly with the number of tokens but rather logarithmically. The ratio plot is therefore more informative for comparing estimator behavior than the raw $\hat{S}$ plot. It shows that the amount of extrapolation relative to the observed vocabulary is broadly similar across languages for a given estimator, while estimator choice still has a large effect.




Figure 14: Median estimated-to-observed richness ratio by language and estimator.

6.7 Estimator agreement and disagreement

Finally, we compare the estimators directly across all spoken corpora.

Estimator	Mean ratio	Median ratio	IQR	Failure rate	Typical behavior
Chao1	2.2673	2.2119	0.7028	0.0	Lower bound via f_1/f_2 ratio
iChao1	2.4760	2.4246	0.7789	0.0	Tighter lower bound using f_3/f_4
ACE	3.9731	3.9421	0.9818	0.1	Heterogeneity-corrected lower bound
Jackknife1	1.5790	1.5998	0.1413	0.0	Delete-one jackknife
Jackknife2	2.0236	2.0624	0.2837	0.0	Delete-two jackknife
Breakaway	13.6129	5.0446	3.8501	50.7	Ratio-regression; unstable for small n

Table 9: Global estimator behavior across all spoken corpora. Ratios are $\hat{S}/S_{obs}$.

In table 9, the global summary shows a clear ordering of estimator aggressiveness. Jackknife1 gives the smallest median ratio, Chao1 and Jackknife2 give moderate ratios, iChao1 is slightly above Chao1, ACE gives larger estimates because it corrects for heterogeneity among rare types, and Breakaway gives the largest median ratio. The very large mean ratio for Breakaway indicates that a small number of successful fits produce extremely large estimates.

Chao1, iChao1, and Jackknife1 are numerically stable on all units. ACE and Jackknife2 fail only rarely. Breakaway fails for about half of the units, meaning its summaries are based only on the subset of corpora where the ratio-regression fit is possible and valid. This makes Breakaway less reliable as a general-purpose estimator for the present data (table 10).

Estimator	Total units	Failed/undefined	Failure rate (%)
Chao1	16374	0	0.0
iChao1	16374	0	0.0
ACE	16374	10	0.1
Jackknife1	16374	0	0.0
Jackknife2	16374	2	0.0
Breakaway	16374	8300	50.7

Table 10: Estimator failure and undefined-output summary for spoken corpora.




Figure 15: Ratio between each estimator and Chao1 across the spoken corpora.

In Figure 15 the ratios relative to Chao1 show how each estimator scales compared with a familiar lower-bound baseline. Values above 1 mean that an estimator is more aggressive than Chao1, while values below 1 mean that it

gives smaller estimates. This representation is easier to compare across corpora of different absolute sizes.

6.8 Summary

The results of this section paint a coherent picture. Across all three languages and all estimators, the observed vocabulary captures only a fraction of the estimated total, ranging between a half and a third, depending on the estimator. The choice of estimator matters far more than the choice of language: median ratios range from around 1.6 for Jackknife1 to over 5 for Breakaway, while for any given estimator the ratios are consistent across English, French, and German. The near-perfect agreement of shape-constrained estimators with Chao1 and iChao1 suggests that the finite-k model is sufficient for this type of data. Breakaway, while being theoretically appealing, fails on half the corpus and should be used with caution for short spoken samples. These findings set up the conclusion in Section 7, where we interpret the cross-language consistency as evidence that the gap between what a speaker says and what they could say is a relatively universal property of spoken language.

7 Conclusion

This thesis set out to answer a simple question: how large is a person’s active spoken vocabulary? By framing spoken word production as a species richness problem, where each distinct word type is a species and each token is an individual, we were able to apply a rich toolkit of nonparametric estimators to transcribed speech data across three languages.

The most striking finding of this work is that, once estimates are normalized by the observed vocabulary size, speakers of English, French, and German exhibit remarkably similar richness ratios across all estimators. Despite substantial differences in corpus size, recording context, and linguistic structure, the ratio $\hat{S}/S_{\text{obs}}$ is broadly consistent across languages for any given estimator. This suggests that the gap between what a speaker says and what they *could* say is a relatively universal feature of spoken language, rather than something specific to a particular language or culture. In contrast, the choice of estimator has a much larger effect on the absolute estimates than the choice of language, with median ratios ranging from around 1.6 for Jackknife1 to over 5 for Breakaway.

The shape-constrained estimators, i.e., the finite k -monotone model of Chee and Wang [8] and the completely monotone estimator of Balabdaoui and Kulagina [2], produce estimates that are in near-perfect agreement with each other, and the AIC-selected degree of monotonicity k^* is moderate across all languages and corpora. This confirms that the data is well described by a moderately flexible shape-constrained model and that the completely monotone limit is not strictly required by these data. Breakaway, on the contrary, fails for roughly half of all units, making it unreliable as a general-purpose estimator for short spoken samples.

These results paint a consistent picture: a typical spoken sample captures only a fraction of a speaker’s active vocabulary, and substantial extrapolation is required regardless of the language or estimator used. Future work could investigate whether similarity in normalized richness ratios across languages persists when corpora are more carefully matched in size, register, and speaker demographics, and whether lemmatization or other linguistic normalization choices materially affect the estimates.

A Parametric mixture models

We give the full probability mass functions for the three parametric Poisson mixture models fitted in Section 4.1. In each case, the conditional model is $X \mid \Lambda = \lambda \sim \text{Poisson}(\lambda)$; the models differ only in the mixing distribution assumed for Λ .

Gamma–Poisson mixture (Negative Binomial). Assume $\Lambda \sim \text{Gamma}(r, r/\mu)$. Integrating Λ yields the negative binomial distribution

$$p(x; r, \mu) = \frac{\Gamma(r+x)}{\Gamma(r) x!} \left(\frac{r}{r+\mu} \right)^r \left(\frac{\mu}{r+\mu} \right)^x, \quad x \in \mathbb{N}_0.$$

This is a classical two-parameter family introduced by Fisher et al. [10] to model species abundances.

LogNormal–Poisson mixture. Assume $\log \Lambda \sim \mathcal{N}(\mu, \sigma^2)$. The marginal distribution has no closed form and is given by

$$p(x; \mu, \sigma) = \int_0^\infty \frac{\lambda^x e^{-\lambda}}{x!} \frac{1}{\lambda \sigma \sqrt{2\pi}} \exp\left(-\frac{(\log \lambda - \mu)^2}{2\sigma^2}\right) d\lambda, \quad x \in \mathbb{N}_0.$$

The integral must be evaluated numerically.

Poisson–Inverse Gaussian mixture. Assume $\Lambda \sim \text{InvGaussian}(\mu, \phi)$. The marginal distribution is

$$p(x; \mu, \phi) = \int_0^\infty \frac{\lambda^x e^{-\lambda}}{x!} dG_{\text{IG}}(\lambda), \quad x \in \mathbb{N}_0,$$

where G_{IG} denotes the inverse Gaussian CDF. The probabilities can be computed exactly from the Taylor coefficients of the probability generating function

$$G(z) = \exp\left(\frac{\phi}{\mu} \left(1 - \sqrt{1 + \frac{2\mu^2(1-z)}{\phi}}\right)\right).$$

B Data Sources

- **British National Corpus (BNC)**
<https://www.english-corpora.org/bnc/>
- **Santa Barbara Corpus of Spoken American English**
<https://www.linguistics.ucsb.edu/research/santa-barbara-corpus-spoken-american-english>
- **Corpus of Contemporary American English (COCA), Spoken Component**
<https://www.english-corpora.org/coca-spoken.asp>
- **Internet Movie Script Database (IMSDb)**
<https://imsdb.com/>

- **CLAIRE Dialog French Dataset**

A collection of transcripts of real French dialogs from various sources, including parliamentary proceedings, interviews, debates, meetings, and free conversations.

<https://huggingface.co/datasets/OpenLLM-France/Claire-Dialogue-French-0.1>

- **Datenbank Gesprochenes Deutsch (DGD)**

A database of transcribed spoken German maintained by the Leibniz Institute for the German Language (IDS Mannheim).

https://dgd.ids-mannheim.de/dgd/pragdb.dgd_extern.welcome

References

- [1] R. Harald Baayen. *Word Frequency Distributions*, volume 18 of *Text, Speech and Language Technology*. Kluwer Academic Publishers, Dordrecht, 2001.
- [2] Fadoua Balabdaoui and Yulia Kulagina. Completely monotone distributions: Mixing, approximation and estimation of number of species. *Computational Statistics and Data Analysis*, 150:107014, 2020.
- [3] Kenneth P. Burnham and W. Scott Overton. Robust estimation of population size when capture probabilities vary among animals. *Ecology*, 60(5):927–936, 1979.
- [4] Anne Chao. Nonparametric estimation of the number of classes in a population. *Scandinavian Journal of Statistics*, 11(4):265–270, 1984.
- [5] Anne Chao and Chun-Huo Chiu. Species richness: Estimation and comparison. *Wiley StatsRef: Statistics Reference Online*, pages 1–26, 2016.
- [6] Anne Chao and Lou Jost. Coverage-based rarefaction and extrapolation: standardizing samples by completeness rather than size. *Ecology*, 93(12):2533–2547, 2012.
- [7] Anne Chao and Shen-Ming Lee. Estimating the number of classes via sample coverage. *Journal of the American Statistical Association*, 87(417):210–217, 1992.
- [8] Chew-Seng Chee and Yong Wang. Nonparametric estimation of species richness using discrete k -monotone distributions. *Computational Statistics and Data Analysis*, 93:107–118, 2016.
- [9] Chun-Huo Chiu, Yi-Ting Wang, Bruno A. Walther, and Anne Chao. An improved nonparametric lower bound of species richness via a modified Good–Turing frequency formula. *Biometrics*, 70(3):671–682, 2014.
- [10] R. A. Fisher, A. S. Corbet, and C. B. Williams. The relation between the number of species and the number of individuals in a random sample of an animal population. *Journal of Animal Ecology*, 12(1):42–58, 1943.
- [11] William A. Gale and Geoffrey Sampson. Good-Turing frequency estimation without tears. *Journal of Quantitative Linguistics*, 2(3):217–237, 1995.
- [12] I. J. Good. The population frequencies of species and the estimation of population parameters. *Biometrika*, 40(3–4):237–264, 1953.
- [13] I. J. Good. Turing’s anticipation of empirical Bayes in connection with the cryptanalysis of the naval Enigma. *Journal of Statistical Computation and Simulation*, 66:101–111, 2000.
- [14] David Joyce. The Poisson process. Lecture notes, Math 217 Probability and Statistics, Clark University, 2014. Fall 2014.

- [15] Claude Lefèvre and Stéphane Loisel. On multiply monotone distributions, continuous or discrete, with applications. *Journal of Applied Probability*, 50(3):827–847, 2013.
- [16] Lalitha Sanathanan. Estimating the size of a multinomial population. *The Annals of Mathematical Statistics*, 43(1):142–152, 1972.
- [17] Johanna Elena Schmitz and Sven Rahmann. A comprehensive review and evaluation of species richness estimation. *Briefings in Bioinformatics*, 26(2):bbaf158, 2025.
- [18] Marvin Spevack. *A Complete and Systematic Concordance to the Works of Shakespeare*, volume 1–6. George Olms, Hildesheim, 1968.
- [19] F. W. Steutel. Note on completely monotone densities. *The Annals of Mathematical Statistics*, 40:1130–1131, 1969.
- [20] Ji-Ping Z. Wang and Bruce G. Lindsay. A penalized nonparametric maximum likelihood approach to species richness estimation. *Journal of the American Statistical Association*, 100(471):942–959, 2005.
- [21] Amy Willis and John Bunge. Estimating diversity via frequency ratios. *Biometrics*, 71(4):1042–1049, 2015.